

# WRITEUP 2

CS 444

OCTOBER 29, 2018

ANTON SYNYTSIA, EYTAN BRODSKY, DAVID JANSEN

	1
<b>Contents</b>	
<b>1 C-Look Elevator Scheduler</b>	<b>2</b>
<b>2 Compiling Qemu with C-Look</b>	<b>3</b>
2.1 Setup . . . . .	3
2.2 Compiling . . . . .	3
<b>3 Starting Qemu</b>	<b>4</b>
<b>4 Testing</b>	<b>4</b>
<b>5 Concurrency 2</b>	<b>5</b>
<b>6 Version Control Log</b>	<b>5</b>
<b>References</b>	<b>8</b>

## 1 C-Look Elevator Scheduler

C-Look elevator reads/writes from the lowest request address to the highest request address without moving back or forth, at any time in between, except for after dispatching the highest requested address, where the elevator resets to the next, lowest, requested address [1]. This allows for the disk head to read/write efficiently, in an expected direction; otherwise, changing directions too much will have performance implications and can physically deteriorate the hard-drive.

In order to implement C-Look, we had to start with an existing scheduler as a template template:

```
cd /scratch/fall2018/group1/linux-yocto-3.19/block/
cp noop-iosched.c clook-iosched.c
```

We then renamed all functions and identifiers C-Look, and modified two functions:

**Request Adder** Instead of appending to the end, like NO-OP did, we insert at a sorted location within the doubly-linked list of requests. Added requests are stored and sorted to the right of head:

HEAD  $\rightarrow$  1  $\rightarrow$  2  $\rightarrow$  7  $\rightarrow$  33  $\rightarrow$  34  $\rightarrow$  ...

Our sort criteria is blk\_rq\_pos.

In order to insert at the right location, we do the following:

- 1) Start at HEAD and traverse forward until:
  - Next points to HEAD, meaning at last request node.
  - Next points to a request node, whose block position is greater than the block position of the request node being added.
- 2) Append request node to currently, iterated node, via list\_add.

```
static void clook_add_request(struct request_queue *q, struct request *rq)
{
    struct clook_data *nd = q->elevator->elevator_data;

    struct list_head *head = &nd->queue;
    struct list_head *item = &rq->queuelist;

    // Find node to append to
    struct list_head *node = head;
    struct list_head *next_node = head->next;

    while (next_node != head) { // while node is not last
        // get request at next node
        struct request *nrq = list_entry(next_node, struct request, queuelist);

        // if request at next node is greater than rq, break out
        if (blk_rq_pos(nrq) > blk_rq_pos(rq)) break;
    }
}
```

```

        // increment
        node = next_node;
        next_node = node->next;
    }

    // Append request item to node
    list_add(item, node);
}

```

**Request Dispatcher** Our modified version of the dispatcher prints out a note in case the request access order is inconsistent, that is if next request is less than previous:

This consistency verification also required a global variable, which we defined at the top:

```
sector_t g_last_seg = 0;
```

The following is a section of code for

## 2 Compiling Qemu with C-Look

### 2.1 Setup

In order to compile Qemu with our scheduler, we had to edit two files within `/scratch/fall2018/group1/linux-yocto-3.19/block/`:

**Makefile** Within the Makefile, we had to add another configuration:

```
obj-$(CONFIG_IOSCHED_CLOOK) += clook-iosched.o
```

**Kconfig.iosched** Within this file, we had to add the following:

```

config IOSCHED_CLOOK
    bool
    default y
    ---help---
    CS444 Group1 Clook scheduler.

```

### 2.2 Compiling

To recompile the scheduler, we ran the following commands:

```

cd /scratch/fall2018/group1/linux-yocto-3.19/
make clean
cp /scratch/files/config-3.19.2-yocto-standard .config
# press y to overwrite
make -j4 all > log_res.txt

```

### 3 Starting Qemu

In order to ease the process of running qemu, as well as, not having to run it in debug environment, we created a `qemu-run.sh` with custom commands:

```
source ./environment-setup-i586-poky-linux
```

```
#qemu-system-i386 -gdb tcp::5501 -S -nographic -kernel bzImage-qemu86.bin -drive
↪ file=core-image-lsb-sdk-qemu86.ext4,if=virtio -enable-kvm -net none -usb -localtime --no-reboot
↪ --append "root=/dev/vda rw console=ttyS0 elevator=clock"
```

```
qemu-system-i386 -gdb tcp::5501 -nographic -kernel
↪ /scratch/fall2018/group1/linux-yocto-3.19/arch/x86/boot/bzImage -drive
↪ file=core-image-lsb-sdk-qemu86.ext4 -enable-kvm -net none -usb -localtime --no-reboot --append
↪ "root=/dev/hda rw console=ttyS0 debug elevator=clock"
```

To run, execute the following commands:

```
# Set path
cd /scratch/fall2018/group1/
# (Optional) Allow execute permission
chmod +x qemu-run.sh
# Run
./qemu-run.sh
```

### 4 Testing

- 1) Create test script.

```
import os
import mmap
import struct
import random

mem = os.open("./big_file", os.O_RDWR)
addr = mmap.mmap(mem, 0x40000000, offset=random.randint(1, 256)*mmap.ALLOCATIONGRANULARITY)

base = 100
i = struct.unpack('I', addr[base:base+4])[0]
print(i)
```

- 2) `cat /dev/urandom >big_file` (create a large file)
- 3) Run the python script on the file. If there are no errors, then the schedule is working.

V	tag	date	commit message	MF	AL	DL
1		2018-10-08	Initial commit	1	1	0
2		2018-10-08	Create .gitignore	1	6	0
Please continue at the next page						



V	tag	date	commit message	MF	AL	DL
34		2018-10-15	Updated work log	5	477	77
35		2018-10-15	Removed thirdparty	1	0	362
36		2018-10-22	Added concurrency2	2	219	0
37		2018-10-23	Setup writeup2 and update comments in concurrency	5	233	4
38		2018-10-27	Unignored Linux Yocto with our sched	7	2	1
39		2018-10-27	Unignoring yocto	1	0	2
40		2018-10-27	yocto...	1	1	0
41		2018-10-27	Added clook scheduler	4	582	2
42		2018-10-27	Setup Writeup2	11	9310	501
43		2018-10-27	Updated makefiles	3	8	9
44		2018-10-27	Making refs work	1	1	1
45		2018-10-27	Added Eyton's scheduler tester	2	16	0
46		2018-10-27	Added David's makefile	2	27	1
47		2018-10-28	Fixed compile errors in CLOOK and modifed Kconfig for inclusion	4	84	6
48		2018-10-28	Added command for starting VM	1	5	0
49		2018-10-29	Made clook print if unexpected access order	2	8	0
50		2018-10-29	Clook Scheduler Writeup	11	165	8
51		2018-10-29	Fixes to README	1	7	7



## References

- [1] “Disk scheduling,” <https://web.archive.org/web/20080606005055/http://www.dcs.ed.ac.uk/home/stg/pub/D/disk.html>.